

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

In [2]:

```
#data input
df = pd.read_csv("../DATA/Ames_Housing_Data.csv")
```


In [3]:

```
df.head()
```

Out[3]:

	PID	MS SubClass	MS Zoning	Lot Frontage	Lot Area	Street	Alley	Lot Shape	Land Contour	Utilities	...
0	526301100	20	RL	141.0	31770	Pave	NaN	IR1	Lvl	AllPub	...
1	526350040	20	RH	80.0	11622	Pave	NaN	Reg	Lvl	AllPub	...
2	526351010	20	RL	81.0	14267	Pave	NaN	IR1	Lvl	AllPub	...
3	526353030	20	RL	93.0	11160	Pave	NaN	Reg	Lvl	AllPub	...
4	527105010	60	RL	74.0	13830	Pave	NaN	IR1	Lvl	AllPub	...

5 rows × 81 columns



In [4]:

```
#correlation between all features and sales price in sorted order
#positive co-relation and value close to 1 means sales price closely depends on that parameter
#in correlation we can see sales price highly depends on Overall Qual
df.corr()['SalePrice'].sort_values()
```

Out[4]:


```
PID -0.246521
Enclosed Porch -0.128787
Kitchen AbvGr -0.119814
Overall Cond -0.101697
MS SubClass -0.085092
Low Qual Fin SF -0.037660
Bsmt Half Bath -0.035835
Yr Sold -0.030569
Misc Val -0.015691
BsmtFin SF 2 0.005891
3Ssn Porch 0.032225
Mo Sold 0.035259
Pool Area 0.068403
Screen Porch 0.112151
Bedroom AbvGr 0.143913
Bsmt Unf SF 0.182855
Lot Area 0.266549
2nd Flr SF 0.269373
Bsmt Full Bath 0.276050
Half Bath 0.285056
Open Porch SF 0.312951
Wood Deck SF 0.327143
Lot Frontage 0.357318
BsmtFin SF 1 0.432914
Fireplaces 0.474558
TotRms AbvGrd 0.495474
Mas Vnr Area 0.508285
Garage Yr Blt 0.526965
Year Remod/Add 0.532974
Full Bath 0.545604
Year Built 0.558426
1st Flr SF 0.621676
Total Bsmt SF 0.632280
Garage Area 0.640401
Garage Cars 0.647877
Gr Liv Area 0.706780
Overall Qual 0.799262
SalePrice 1.000000
Name: SalePrice, dtype: float64
```

In [5]:

```
#since sales price highly depends on overall quality  
sns.scatterplot(x='Overall Qual',y='SalePrice',data=df)  
#We notice higher quality higher is sales price but their are few outlier points which we n
```

Out[5]:

<AxesSubplot:xlabel='Overall Qual', ylabel='SalePrice'>



In [6]:

```
#from above we notice there are 3 houses which have quality between 8-10 but are selling ve  
#Therefore they are doubtfull point
```

In [7]:

```
df[(df['Overall Qual']>8) & (df['SalePrice']<200000)]
```

Out[7]:

	PID	MS SubClass	MS Zoning	Lot Frontage	Lot Area	Street	Alley	Lot Shape	Land Contour	Utilities
1182	533350090	60	RL	NaN	24572	Pave	NaN	IR1	Lvl	AllPub
1498	908154235	60	RL	313.0	63887	Pave	NaN	IR3	Bnk	AllPub
2180	908154195	20	RL	128.0	39290	Pave	NaN	IR1	Bnk	AllPub
2181	908154205	60	RL	130.0	40094	Pave	NaN	IR1	Bnk	AllPub

4 rows × 81 columns

In [8]:


```
#The four suspicious rows are displayed above
```

In [9]:

```
#Gr Liv Area is also highly correlated to sales price
sns.scatterplot(x='Gr Liv Area',y='SalePrice',data=df)
```

Out[9]:

```
<AxesSubplot:xlabel='Gr Liv Area', ylabel='SalePrice'>
```



In [10]:

```
#We again notice that the above mentoined three houses here again show weird behaviour
#The general trend is higher the Gr Liv Area more is the SalesPrice
#But those three houses Have high Gr Liv Area But small selling price
```

In [11]:

```
df[(df['Gr Liv Area']>4000) & (df['SalePrice']<400000)]
```

Out[11]:

	PID	MS SubClass	MS Zoning	Lot Frontage	Lot Area	Street	Alley	Lot Shape	Land Contour	Utilities
1498	908154235	60	RL	313.0	63887	Pave	NaN	IR3	Bnk	AllPu
2180	908154195	20	RL	128.0	39290	Pave	NaN	IR1	Bnk	AllPu
2181	908154205	60	RL	130.0	40094	Pave	NaN	IR1	Bnk	AllPu

3 rows × 11 columns

In [12]:

```
#We get the three rows with high Gr Liv Area but low sales price. These three rows match th  
#There are outliers which should be removed
```

In [13]:

```
df[(df['Gr Liv Area']>4000) & (df['SalePrice']<400000)].index
```

Out[13]:

```
Int64Index([1498, 2180, 2181], dtype='int64')
```

In [14]:

```
ind_drop = df[(df['Gr Liv Area']>4000) & (df['SalePrice']<400000)].index
```

In [15]:


```
#Those three rows dropped  
df = df.drop(ind_drop,axis=0)
```

In [16]:

```
sns.scatterplot(x='Gr Liv Area',y='SalePrice',data=df)
```

Out[16]:

<AxesSubplot:xlabel='Gr Liv Area', ylabel='SalePrice'>



In [17]:


```
#Now the rest of the data set seems to follow the general trend
```

In [18]:

```
sns.scatterplot(x='Overall Qual',y='SalePrice',data=df)
```

Out[18]:

<AxesSubplot:xlabel='Overall Qual', ylabel='SalePrice'>



In [19]:

```
#Rest of data seems to follow the general trend
```

In []:

In [20]:

```
#WE DEALT WITH OUTLIERS NOW WE DEAL WITH MISSING DATA
```

In [21]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2927 entries, 0 to 2929
Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   PID                   2927 non-null  int64
1   MS SubClass           2927 non-null  int64
2   MS Zoning              2927 non-null  object
3   Lot Frontage          2437 non-null  float64
4   Lot Area              2927 non-null  int64
5   Street                2927 non-null  object
6   Alley                 198 non-null   object
7   Lot Shape             2927 non-null  object
8   Land Contour          2927 non-null  object
9   Utilities             2927 non-null  object
10  Lot Config            2927 non-null  object
11  Land Slope            2927 non-null  object
12  Neighborhood          2927 non-null  object
13  Condition 1           2927 non-null  object
14  Condition 2           2927 non-null  object
15  Bldg Type             2927 non-null  object
16  House Style           2927 non-null  object
17  Overall Qual          2927 non-null  int64
18  Overall Cond          2927 non-null  int64
19  Year Built            2927 non-null  int64
20  Year Remod/Add        2927 non-null  int64
21  Roof Style            2927 non-null  object
22  Roof Matl             2927 non-null  object
23  Exterior 1st          2927 non-null  object
24  Exterior 2nd          2927 non-null  object
25  Mas Vnr Type          2904 non-null  object
26  Mas Vnr Area          2904 non-null  float64
27  Exter Qual            2927 non-null  object
28  Exter Cond            2927 non-null  object
29  Foundation            2927 non-null  object
30  Bsmt Qual             2847 non-null  object
31  Bsmt Cond             2847 non-null  object
32  Bsmt Exposure         2844 non-null  object
33  BsmtFin Type 1        2847 non-null  object
34  BsmtFin SF 1          2926 non-null  float64
35  BsmtFin Type 2        2846 non-null  object
36  BsmtFin SF 2          2926 non-null  float64
37  Bsmt Unf SF           2926 non-null  float64
38  Total Bsmt SF         2926 non-null  float64
39  Heating               2927 non-null  object
40  Heating QC            2927 non-null  object
41  Central Air           2927 non-null  object
42  Electrical            2926 non-null  object
43  1st Flr SF            2927 non-null  int64
44  2nd Flr SF            2927 non-null  int64
45  Low Qual Fin SF       2927 non-null  int64
46  Gr Liv Area           2927 non-null  int64
47  Bsmt Full Bath        2925 non-null  float64
48  Bsmt Half Bath        2925 non-null  float64
49  Full Bath             2927 non-null  int64
50  Half Bath             2927 non-null  int64
51  Bedroom AbvGr         2927 non-null  int64
```



```
52 Kitchen AbvGr 2927 non-null int64
53 Kitchen Qual 2927 non-null object
54 TotRms AbvGrd 2927 non-null int64
55 Functional 2927 non-null object
56 Fireplaces 2927 non-null int64
57 Fireplace Qu 1505 non-null object
58 Garage Type 2770 non-null object
59 Garage Yr Blt 2768 non-null float64
60 Garage Finish 2768 non-null object
61 Garage Cars 2926 non-null float64
62 Garage Area 2926 non-null float64
63 Garage Qual 2768 non-null object
64 Garage Cond 2768 non-null object
65 Paved Drive 2927 non-null object
66 Wood Deck SF 2927 non-null int64
67 Open Porch SF 2927 non-null int64
68 Enclosed Porch 2927 non-null int64
69 3Ssn Porch 2927 non-null int64
70 Screen Porch 2927 non-null int64
71 Pool Area 2927 non-null int64
72 Pool QC 12 non-null object
73 Fence 572 non-null object
74 Misc Feature 105 non-null object
75 Misc Val 2927 non-null int64
76 Mo Sold 2927 non-null int64
77 Yr Sold 2927 non-null int64
78 Sale Type 2927 non-null object
79 Sale Condition 2927 non-null object
80 SalePrice 2927 non-null int64
```

dtypes: float64(11), int64(27), object(43)

memory usage: 1.8+ MB

In [22]:

```
#We Notice that there for few features there are missing values
```

In [23]:

```
df = df.drop('PID',axis=1)
```

In [24]:

```
df.head()
```

Out[24]:

	MS SubClass	MS Zoning	Lot Frontage	Lot Area	Street	Alley	Lot Shape	Land Contour	Utilities	Lot Config	...	Pot Area
0	20	RL	141.0	31770	Pave	NaN	IR1	Lvl	AllPub	Corner	...	
1	20	RH	80.0	11622	Pave	NaN	Reg	Lvl	AllPub	Inside	...	
2	20	RL	81.0	14267	Pave	NaN	IR1	Lvl	AllPub	Corner	...	
3	20	RL	93.0	11160	Pave	NaN	Reg	Lvl	AllPub	Corner	...	
4	60	RL	74.0	13830	Pave	NaN	IR1	Lvl	AllPub	Inside	...	

5 rows × 80 columns

In [25]:

```
df.isnull().sum()
```

Out[25]:

```
MS SubClass      0
MS Zoning        0
Lot Frontage    490
Lot Area         0
Street          0
...
Mo Sold         0
Yr Sold         0
Sale Type       0
Sale Condition  0
SalePrice       0
Length: 80, dtype: int64
```

In [26]:

```
#True treated as 0 and false as 1 therefore we get sum of how many rows for each feature are 0
#there are 80 features therefore we cannot see all
```

In [27]:

```
100* df.isnull().sum() / len(df)
```

Out[27]:

```
MS SubClass      0.00000
MS Zoning        0.00000
Lot Frontage     16.74069
Lot Area         0.00000
Street           0.00000
...
Mo Sold          0.00000
Yr Sold          0.00000
Sale Type        0.00000
Sale Condition   0.00000
SalePrice        0.00000
Length: 80, dtype: float64
```

In [28]:

```
#We now get what percentage of data is missing which would help us to evaluate better
```

In [29]:

```
def percent_missing(df):
    percent_nan = 100* df.isnull().sum() / len(df)
    percent_nan = percent_nan[percent_nan>0].sort_values()
    return percent_nan
```

In [30]:

```
percent_nan = percent_missing(df)
```

In [31]:

```
percent_nan
```

Out[31]:


```
Electrical      0.034165
Garage Area     0.034165
Total Bsmt SF   0.034165
Bsmt Unf SF     0.034165
BsmtFin SF 1    0.034165
BsmtFin SF 2    0.034165
Garage Cars     0.034165
Bsmt Full Bath  0.068329
Bsmt Half Bath  0.068329
Mas Vnr Area    0.785787
Mas Vnr Type    0.785787
BsmtFin Type 1  2.733174
Bsmt Qual       2.733174
Bsmt Cond       2.733174
BsmtFin Type 2  2.767339
Bsmt Exposure   2.835668
Garage Type     5.363854
Garage Finish   5.432183
Garage Qual     5.432183
Garage Cond     5.432183
Garage Yr Blt   5.432183
Lot Frontage    16.740690
Fireplace Qu    48.582166
Fence           80.457807
Alley           93.235395
Misc Feature    96.412709
Pool QC         99.590024
dtype: float64
```

In [32]:

```
#We get percentage of missing data in sorted manner using a function call
```

In [33]:

```
sns.barplot(x=percent_nan.index,y=percent_nan)
plt.xticks(rotation=90);
```



In [34]:

```
#graphical representation of missing data
```


In [35]:

```
sns.barplot(x=percent_nan.index,y=percent_nan)
plt.xticks(rotation=90);
```

```
# Set 1% Threshold
plt.ylim(0,1)
```

Out[35]:

(0.0, 1.0)



In [36]:

```
#I would not mind dropping rows with 1% missing data
```

In [37]:

```
percent_nan[percent_nan < 1]
```

Out[37]:

```
Electrical      0.034165
Garage Area     0.034165
Total Bsmt SF   0.034165
Bsmt Unf SF     0.034165
BsmtFin SF 1    0.034165
BsmtFin SF 2    0.034165
Garage Cars     0.034165
Bsmt Full Bath  0.068329
Bsmt Half Bath  0.068329
Mas Vnr Area    0.785787
Mas Vnr Type    0.785787
dtype: float64
```

In [38]:

```
#Features with less than 1% missing data
```

In [39]:

```
100/len(df)
```

Out[39]:

```
0.0341646737273659
```

In [40]:

```
#the above calculation tells that features like Electrical, Garage Area, Bsmt Unf SF etc ha  
#is missing data
```

In [41]:

```
df = df.dropna(axis = 0, subset = ['Electrical', 'Garage Area'])
```

In [42]:

```
percent_nan = percent_missing(df)
```

In [43]:

```
percent_nan[percent_nan < 1]
```

Out[43]:

```
Bsmt Unf SF      0.034188
Total Bsmt SF    0.034188
BsmtFin SF 2     0.034188
BsmtFin SF 1     0.034188
Bsmt Full Bath   0.068376
Bsmt Half Bath   0.068376
Mas Vnr Type     0.786325
Mas Vnr Area     0.786325
dtype: float64
```

In [44]:

```
#We notice by dropping electrical and garage area we also dropped many other features which
#same row as above two features missing
```

In [45]:

```
df[df['Bsmt Half Bath'].isnull()]
```

Out[45]:

	MS SubClass	MS Zoning	Lot Frontage	Lot Area	Street	Alley	Lot Shape	Land Contour	Utilities	Lot Config	...
1341	20	RM	99.0	5940	Pave	NaN	IR1	Lvl	AllPub	FR3	...
1497	20	RL	123.0	47007	Pave	NaN	IR1	Lvl	AllPub	Inside	...

2 rows × 80 columns



In [46]:

```
df[df['Bsmt Full Bath'].isnull()]
```

Out[46]:

	MS SubClass	MS Zoning	Lot Frontage	Lot Area	Street	Alley	Lot Shape	Land Contour	Utilities	Lot Config	...
1341	20	RM	99.0	5940	Pave	NaN	IR1	Lvl	AllPub	FR3	...
1497	20	RL	123.0	47007	Pave	NaN	IR1	Lvl	AllPub	Inside	...

2 rows × 80 columns



In [47]:

```
#We notice that the same two rows are missing information for full bath and half bath
```

In [48]:

```
with open('../DATA/Ames_Housing_Feature_Description.txt','r') as f:
    print(f.read())
```

MSSubClass: Identifies the type of dwelling involved in the sale.

```

20      1-STORY 1946 & NEWER ALL STYLES
30      1-STORY 1945 & OLDER
40      1-STORY W/FINISHED ATTIC ALL AGES
45      1-1/2 STORY - UNFINISHED ALL AGES
50      1-1/2 STORY FINISHED ALL AGES
60      2-STORY 1946 & NEWER
70      2-STORY 1945 & OLDER
75      2-1/2 STORY ALL AGES
80      SPLIT OR MULTI-LEVEL
85      SPLIT FOYER
90      DUPLEX - ALL STYLES AND AGES
120     1-STORY PUD (Planned Unit Development) - 1946 & NEWER
150     1-1/2 STORY PUD - ALL AGES
160     2-STORY PUD - 1946 & NEWER
180     PUD - MULTILEVEL - INCL SPLIT LEV/FOYER
190     2 FAMILY CONVERSION - ALL STYLES AND AGES
```

MZoning: Identifies the general zoning classification of the sale.

In [49]:

```
#In data description we see NA for various basement parameters means there is no basement i
#Therefore instead of dropping them for numerical basement values we can fill them with zer
#And for string basement values we can fill them with none thus telling that there are no b
```

In [50]:

```
#BSMT numeric coloumn --> 0
bsmt_num_cols = ['BsmtFin SF 1', 'BsmtFin SF 2', 'Bsmt Unf SF', 'Total Bsmt SF', 'Bsmt Full
df[bsmt_num_cols] = df[bsmt_num_cols].fillna(0)
```

In [51]:


```
#BSMT string colom --> NONE
bsmt_str_cols = ['Bsmt Qual', 'Bsmt Cond', 'Bsmt Exposure', 'BsmtFin Type 1', 'BsmtFin Typ
df[bsmt_str_cols] = df[bsmt_str_cols].fillna('None')
```

In [52]:

```
percent_nan = percent_missing(df)
```


In [53]:

```
sns.barplot(x=percent_nan.index,y=percent_nan)
plt.xticks(rotation=90);
```



In [54]:

```
#Now again in case of Mas VNR Type and Area when we refer dataframe detail we find there no  
#Therefore repeat same above process for them as well
```

In [55]:

```
df['Mas Vnr Type'] = df['Mas Vnr Type'].fillna('None')
```

In [56]:


```
df['Mas Vnr Area'] = df['Mas Vnr Area'].fillna(0)
```

In [57]:

```
percent_nan = percent_missing(df)
```

In [58]:

```
sns.barplot(x=percent_nan.index,y=percent_nan)
plt.xticks(rotation=90);
```



In [59]:

#Now we are above the 1% threshold range for dropping rows, therefore we need to consider c

In [60]:

#Now again for garage columns NA n=means no garage exit therefore we repeat same above proce

In [61]:

```
gar_str_cols = ['Garage Type', 'Garage Finish', 'Garage Qual', 'Garage Cond']
df[gar_str_cols] = df[gar_str_cols].fillna('None')
```

In [62]:


```
df['Garage Yr Blt'] = df['Garage Yr Blt'].fillna(0)
```

In [63]:

```
percent_nan = percent_missing(df)
```

In [64]:

```
sns.barplot(x=percent_nan.index,y=percent_nan)
plt.xticks(rotation=90);
```



In [65]:

```
#Fence Alley Misc Feature Pool QC are missing more than 80% of data therefore it is better
```

In [66]:


```
df = df.drop(['Pool QC','Misc Feature','Alley','Fence'],axis=1)
```

In [67]:

```
percent_nan = percent_missing(df)
```

In [68]:

```
sns.barplot(x=percent_nan.index,y=percent_nan)
plt.xticks(rotation=90);
```



In [69]:

```
#In remaining two features we can neither drop rows because too many rows are missing
#nore we can drop coloms because not that much data is missing
```

In [70]:

```
df['Fireplace Qu'].value_counts()
```

Out[70]:

```
Gd      741
TA      600
Fa       75
Po       46
Ex       43
Name: Fireplace Qu, dtype: int64
```

In [71]:

```
#We notice it is a string colom therefore we just fill None in missing values
```

In [72]:


```
df['Fireplace Qu'] = df['Fireplace Qu'].fillna("None")
```

In [73]:

```
percent_nan = percent_missing(df)
```

In [74]:

```
sns.barplot(x=percent_nan.index,y=percent_nan)
plt.xticks(rotation=90);
```



In [75]:

```
df['Lot Frontage']
```

Out[75]:

```
0      141.0
1       80.0
2       81.0
3       93.0
4       74.0
...
2925    37.0
2926     NaN
2927    62.0
2928    77.0
2929    74.0
Name: Lot Frontage, Length: 2925, dtype: float64
```

In [76]:

```
with open('../DATA/Ames_Housing_Feature_Description.txt','r') as f:
    print(f.read())
```

MSSubClass: Identifies the type of dwelling involved in the sale.

```

20      1-STORY 1946 & NEWER ALL STYLES
30      1-STORY 1945 & OLDER
40      1-STORY W/FINISHED ATTIC ALL AGES
45      1-1/2 STORY - UNFINISHED ALL AGES
50      1-1/2 STORY FINISHED ALL AGES
60      2-STORY 1946 & NEWER
70      2-STORY 1945 & OLDER
75      2-1/2 STORY ALL AGES
80      SPLIT OR MULTI-LEVEL
85      SPLIT FOYER
90      DUPLEX - ALL STYLES AND AGES
120     1-STORY PUD (Planned Unit Development) - 1946 & NEWER
150     1-1/2 STORY PUD - ALL AGES
160     2-STORY PUD - 1946 & NEWER
180     PUD - MULTILEVEL - INCL SPLIT LEV/FOYER
190     2 FAMILY CONVERSION - ALL STYLES AND AGES
```

MSTZ = Major Street Ties, a variable that indicates whether the house is located on a major street.

In [77]:

#Thus we see Lot Frontage is Numeric data and it is Linear feet of street connected to prop

In [78]:

```

# Neighborhood: Physical Locations within Ames city Limits

# LotFrontage: Linear feet of street connected to property

# We will operate under the assumption that the Lot Frontage is related to what neighborhood
```

In [79]:

```
df['Neighborhood'].unique()
```

Out[79]:

```
array(['NAmes', 'Gilbert', 'StoneBr', 'NWAmes', 'Somerst', 'BrDale',
      'NPkVill', 'NridgHt', 'Blmngtn', 'NoRidge', 'SawyerW', 'Sawyer',
      'Greens', 'BrkSide', 'OldTown', 'IDOTRR', 'ClearCr', 'SWISU',
      'Edwards', 'CollgCr', 'Crawfor', 'Blueste', 'Mitchel', 'Timber',
      'MeadowV', 'Veenker', 'GrnHill', 'Landmrk'], dtype=object)
```

In [80]:


```
#All neighbourhoods
```

In [81]:

```
plt.figure(figsize=(8,12))
sns.boxplot(x='Lot Frontage',y='Neighborhood',data=df,orient='h')
```

Out[81]:

```
<AxesSubplot:xlabel='Lot Frontage', ylabel='Neighborhood'>
```



In [82]:

```
df.groupby('Neighborhood')['Lot Frontage']
```

Out[82]:

```
<pandas.core.groupby.generic.SeriesGroupBy object at 0x0000016520CBDFa0>
```

In [83]:

```
df.groupby('Neighborhood')['Lot Frontage'].mean()
```

Out[83]:

```
Neighborhood
Blmngtn      46.900000
Blueste      27.300000
BrDale        21.500000
BrkSide       55.789474
ClearCr       88.150000
CollgCr       71.336364
Crawfor       69.951807
Edwards       64.794286
Gilbert       74.207207
Greens        41.000000
GrnHill        NaN
IDOTRR        62.383721
Landmrk        NaN
MeadowV       25.606061
Mitchel       75.144444
NAmes         75.210667
NPkVill        28.142857
NWAmes        81.517647
NoRidge       91.629630
NridgHt       84.184049
OldTown       61.777293
SWISU         59.068182
Sawyer        74.551020
SawyerW       70.669811
Somerst       64.549383
StoneBr       62.173913
Timber        81.303571
Veenker       72.000000
Name: Lot Frontage, dtype: float64
```

In [84]:

```
#Therefore now we have average lot frontage value for each neighbourhood  
#for missing data we fill it with average value of that neighbourhood
```


In [85]:

```
df.groupby('Neighborhood')['Lot Frontage'].transform(lambda val: val.fillna(val.mean()))
```

Out[85]:

```
0      141.000000
1       80.000000
2       81.000000
3       93.000000
4       74.000000
...
2925    37.000000
2926    75.144444
2927    62.000000
2928    77.000000
2929    74.000000
```

Name: Lot Frontage, Length: 2925, dtype: float64

In [86]:

```
df['Lot Frontage'] = df.groupby('Neighborhood')['Lot Frontage'].transform(lambda val: val.f
```

In [87]:


```
#We filled the missing Lot Frontage values
```

In [88]:

```
percent_nan = percent_missing(df)
```

In [89]:

```
sns.barplot(x=percent_nan.index,y=percent_nan)
plt.xticks(rotation=90);
```



In [90]:

```
df['Lot Frontage'] = df['Lot Frontage'].fillna(0)
```

In [91]:

```
percent_nan = percent_missing(df)
```

Dealing with categorical data

In [92]:

```
# Convert to String
df['MS SubClass'] = df['MS SubClass'].apply(str)
```

In [93]:

```
df.select_dtypes(include='object')
```

Out[93]:

	MS SubClass	MS Zoning	Street	Lot Shape	Land Contour	Utilities	Lot Config	Land Slope	Neighborhood	Con
0	20	RL	Pave	IR1	Lvl	AllPub	Corner	Gtl	NAmes	
1	20	RH	Pave	Reg	Lvl	AllPub	Inside	Gtl	NAmes	
2	20	RL	Pave	IR1	Lvl	AllPub	Corner	Gtl	NAmes	
3	20	RL	Pave	Reg	Lvl	AllPub	Corner	Gtl	NAmes	
4	60	RL	Pave	IR1	Lvl	AllPub	Inside	Gtl	Gilbert	
...	
2925	80	RL	Pave	IR1	Lvl	AllPub	CulDSac	Gtl	Mitchel	
2926	20	RL	Pave	IR1	Low	AllPub	Inside	Mod	Mitchel	
2927	85	RL	Pave	Reg	Lvl	AllPub	Inside	Gtl	Mitchel	
2928	20	RL	Pave	Reg	Lvl	AllPub	Inside	Mod	Mitchel	
2929	60	RL	Pave	Reg	Lvl	AllPub	Inside	Mod	Mitchel	

2925 rows × 40 columns



In [94]:

```
df_nums = df.select_dtypes(exclude='object')
df_objs = df.select_dtypes(include='object')
```

In [95]:

df_nums.info()

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2925 entries, 0 to 2929
Data columns (total 36 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Lot Frontage           2925 non-null   float64
1   Lot Area               2925 non-null   int64
2   Overall Qual           2925 non-null   int64
3   Overall Cond           2925 non-null   int64
4   Year Built             2925 non-null   int64
5   Year Remod/Add         2925 non-null   int64
6   Mas Vnr Area           2925 non-null   float64
7   BsmtFin SF 1           2925 non-null   float64
8   BsmtFin SF 2           2925 non-null   float64
9   Bsmt Unf SF            2925 non-null   float64
10  Total Bsmt SF          2925 non-null   float64
11  1st Flr SF             2925 non-null   int64
12  2nd Flr SF             2925 non-null   int64
13  Low Qual Fin SF        2925 non-null   int64
14  Gr Liv Area            2925 non-null   int64
15  Bsmt Full Bath         2925 non-null   float64
16  Bsmt Half Bath         2925 non-null   float64
17  Full Bath              2925 non-null   int64
18  Half Bath              2925 non-null   int64
19  Bedroom AbvGr          2925 non-null   int64
20  Kitchen AbvGr          2925 non-null   int64
21  TotRms AbvGrd          2925 non-null   int64
22  Fireplaces             2925 non-null   int64
23  Garage Yr Blt          2925 non-null   float64
24  Garage Cars            2925 non-null   float64
25  Garage Area            2925 non-null   float64
26  Wood Deck SF           2925 non-null   int64
27  Open Porch SF          2925 non-null   int64
28  Enclosed Porch         2925 non-null   int64
29  3Ssn Porch             2925 non-null   int64
30  Screen Porch           2925 non-null   int64
31  Pool Area              2925 non-null   int64
32  Misc Val               2925 non-null   int64
33  Mo Sold                2925 non-null   int64
34  Yr Sold                2925 non-null   int64
35  SalePrice              2925 non-null   int64
dtypes: float64(11), int64(25)
memory usage: 910.0 KB

```

In [96]:

df_objs.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2925 entries, 0 to 2929
Data columns (total 40 columns):
#   Column                Non-Null Count  Dtype
---  -
0   MS SubClass            2925 non-null   object
1   MS Zoning              2925 non-null   object
2   Street                2925 non-null   object
3   Lot Shape              2925 non-null   object
4   Land Contour           2925 non-null   object
5   Utilities              2925 non-null   object
6   Lot Config             2925 non-null   object
7   Land Slope             2925 non-null   object
8   Neighborhood           2925 non-null   object
9   Condition 1           2925 non-null   object
10  Condition 2            2925 non-null   object
11  Bldg Type              2925 non-null   object
12  House Style            2925 non-null   object
13  Roof Style             2925 non-null   object
14  Roof Matl              2925 non-null   object
15  Exterior 1st           2925 non-null   object
16  Exterior 2nd           2925 non-null   object
17  Mas Vnr Type           2925 non-null   object
18  Exter Qual             2925 non-null   object
19  Exter Cond             2925 non-null   object
20  Foundation             2925 non-null   object
21  Bsmt Qual             2925 non-null   object
22  Bsmt Cond             2925 non-null   object
23  Bsmt Exposure          2925 non-null   object
24  BsmtFin Type 1         2925 non-null   object
25  BsmtFin Type 2         2925 non-null   object
26  Heating                2925 non-null   object
27  Heating QC             2925 non-null   object
28  Central Air            2925 non-null   object
29  Electrical             2925 non-null   object
30  Kitchen Qual           2925 non-null   object
31  Functional             2925 non-null   object
32  Fireplace Qu           2925 non-null   object
33  Garage Type            2925 non-null   object
34  Garage Finish          2925 non-null   object
35  Garage Qual            2925 non-null   object
36  Garage Cond            2925 non-null   object
37  Paved Drive            2925 non-null   object
38  Sale Type              2925 non-null   object
39  Sale Condition         2925 non-null   object
dtypes: object(40)
memory usage: 1001.5+ KB
```

Converting

In [97]:

```
df_objs = pd.get_dummies(df_objs, drop_first=True)
```

In [98]:

```
final_df = pd.concat([df_nums,df_objs],axis=1)
```

In [99]:

```
final_df
```

Out[99]:

	Lot Frontage	Lot Area	Overall Qual	Overall Cond	Year Built	Year Remod/Add	Mas Vnr Area	BsmtFin SF 1	BsmtFin SF 2	Bsmt Unf SF
0	141.000000	31770	6	5	1960	1960	112.0	639.0	0.0	441.0
1	80.000000	11622	5	6	1961	1961	0.0	468.0	144.0	270.0
2	81.000000	14267	6	6	1958	1958	108.0	923.0	0.0	406.0
3	93.000000	11160	7	5	1968	1968	0.0	1065.0	0.0	1045.0
4	74.000000	13830	5	5	1997	1998	0.0	791.0	0.0	137.0
...
2925	37.000000	7937	6	6	1984	1984	0.0	819.0	0.0	184.0
2926	75.144444	8885	5	5	1983	1983	0.0	301.0	324.0	239.0
2927	62.000000	10441	5	5	1992	1992	0.0	337.0	0.0	575.0
2928	77.000000	10010	5	5	1974	1975	0.0	1071.0	123.0	195.0
2929	74.000000	9627	7	5	1993	1994	94.0	758.0	0.0	238.0

2925 rows × 274 columns

In [100]:

```
final_df.corr()['SalePrice'].sort_values()
```

Out[100]:

```
Exter Qual_TA      -0.591459
Kitchen Qual_TA    -0.527461
Fireplace Qu_None  -0.481740
Bsmt Qual_TA       -0.453022
Garage Finish_Unf  -0.422363
...
Garage Cars         0.648488
Total Bsmt SF       0.660983
Gr Liv Area         0.727279
Overall Qual        0.802637
SalePrice           1.000000
Name: SalePrice, Length: 274, dtype: float64
```

Data is cleaned and Now we create the Machine Learning model

We Use Linear Regression Model

In [108]:

```
df = pd.read_csv("../DATA/AMES_Final_DF.csv")
```

In [109]:

```
df.head()
```

Out[109]:

	Lot Frontage	Lot Area	Overall Qual	Overall Cond	Year Built	Year Remod/Add	Mas Vnr Area	BsmtFin SF 1	BsmtFin SF 2	Bsmt Unf SF	...	
0	141.0	31770	6	5	1960	1960	112.0	639.0	0.0	441.0	...	
1	80.0	11622	5	6	1961	1961	0.0	468.0	144.0	270.0	...	
2	81.0	14267	6	6	1958	1958	108.0	923.0	0.0	406.0	...	
3	93.0	11160	7	5	1968	1968	0.0	1065.0	0.0	1045.0	...	
4	74.0	13830	5	5	1997	1998	0.0	791.0	0.0	137.0	...	

5 rows × 274 columns

In [110]:

```
#We are trying to predict sales price label. Therefore we take sales price as y and rest fe
```

In [111]:

```
X = df.drop('SalePrice',axis=1)
```

In [112]:

```
y = df['SalePrice']
```

In [113]:

```
#We would now be using train_test_split from sklearn
```

In [114]:

```
from sklearn.model_selection import train_test_split
```

In [115]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.10, random_state=101)
```

In [116]:

```
#We would now scale the data
```

In [117]:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaled_X_train = scaler.fit_transform(X_train)
scaled_X_test = scaler.transform(X_test)
```

In [118]:

#We use the Elastic model which is a combination of both Lasso and Ridge

In [119]:

```
from sklearn.linear_model import ElasticNet
```

In [120]:

```
base_elastic_model = ElasticNet()
```

In [121]:

*#The Elastic Net model has two main parameters, alpha and the L1 ratio.
#Therefore we create grid of these two to choose most suitable value*

In [122]:

```
param_grid = {'alpha':[0.1,1,5,10,50,100],
              'l1_ratio': [.1, .5, .7, .9, .95, .99, 1]}
```

In [123]:

```
from sklearn.model_selection import GridSearchCV
```

In [124]:

```
grid_model = GridSearchCV(estimator=base_elastic_model,
                          param_grid=param_grid,
                          scoring='neg_mean_squared_error',
                          cv=5,
                          verbose=1)
```

In [125]:

```
grid_model.fit(scaled_X_train,y_train)
```

530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 139422008252.62378, tolerance: 1355206692.5276783

```
model = cd_fast.enet_coordinate_descent(
c:\python39\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
530: ConvergenceWarning: Objective did not converge. You might want to inc
rease the number of iterations. Duality gap: 165405536738.32166, toleranc
e: 1307913805.6588454
```

```
model = cd_fast.enet_coordinate_descent(
c:\python39\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
530: ConvergenceWarning: Objective did not converge. You might want to inc
rease the number of iterations. Duality gap: 132401459345.90894, toleranc
e: 1415056940.0061066
```

```
model = cd_fast.enet_coordinate_descent(
c:\python39\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:
530: ConvergenceWarning: Objective did not converge. You might want to inc
rease the number of iterations. Duality gap: 198623915721.08862, toleranc
e: 1438198040.0882876
```

```
model = cd_fast.enet_coordinate_descent(
c:\nvthon39\lib\site-packages\sklearn\linear model\_coordinate_descent.nv:
```

In [126]:

```
grid_model.best_params_
```

Out[126]:

```
{'alpha': 100, 'l1_ratio': 1}
```

In [127]:

```
y_pred = grid_model.predict(scaled_X_test)
```

In [128]:

```
from sklearn.metrics import mean_absolute_error,mean_squared_error
```

In [129]:

```
mean_absolute_error(y_test,y_pred)
```

Out[129]:

```
14195.354900562173
```

In [130]:

```
np.sqrt(mean_squared_error(y_test,y_pred))
```

Out[130]:

```
20558.508566893164
```


In [131]:

```
np.mean(df['SalePrice'])
```

Out[131]:

180815.53743589742

In [132]:

```
import pickle
```

In [133]:

```
saved_model = pickle.dumps(grid_model)
```

In [134]:

```
#Model is saved
```

In []: